

FROM Expensive Queries TO Smart Caching A DBA's Guide



by Javier Zon



DISCLOSURE



Who's this guy?

- SysAdmin and MySQL DBA
- Founder of ScaleDB
- @jtomaszon
- Taekwondo Black Belt




PROBLEM


Application
Connection
Pooling




Application Connection Pooling



Application Connection Pooling



Application Connection Pooling



PROBLEM

Object
Relational
Mapping "Magic"



Object Relational Mapping "Magic"

```
purchases = Purchase.uncscoped.where(status:"paid").  
  where("original_amount_currency = ?", "USD").  
  joins(:products, :account).  
  group("purchases.account_id").  
pluck("purchases.account_id",  
      "accounts.name",  
      "count(purchases.id)",  
      "sum(original_amount_cents)"  
);
```

Object Relational Mapping "Magic"

```
SELECT
    purchases.account_id,
    accounts.name,
    count(purchases.id),
    sum(original_amount_cents)
FROM `purchases` pu
INNER JOIN `product_purchases` pp ON `pp`.`purchase_id` = `pu`.`id`
INNER JOIN `products` p ON `p`.`id` = `pp`.`product_id`
INNER JOIN `accounts` a ON `a`.`id` = `pu`.`account_id`
WHERE `pu`.`status` = 'paid'
    AND (original_amount_currency = 'USD')
GROUP BY purchases.account_id;
```

PROBLEM

Manual Code for Cache Invalidation

*There are only two hard things in Computer Science:
naming things and **cache invalidation**.*

– Phil Karlton



SOLUTION

- ~~Application Connection Pooling~~ > Smarter Connection Pooling
- ~~Object Relational "Magic"~~ > Visibility into Queries
- ~~Cache invalidation~~ > Cache that stays Valid automatically



SOLUTION

- ProxySQL for Connection Pooling
 - Multiplex DB Connections
 - Routing to different Backends



SOLUTION

```
mysql> select count(*) from stats_mysql_processlist;
```

```
+-----+
| count(*) |
+-----+
| 7415     |
+-----+
```


```
1 row in set (0.27 sec)
```

status	ConnUsed	ConnFree	ConnOK	ConnERR	MaxConnUsed	Queries
ONLINE	0	10	0	0	6	4408
ONLINE	0	0	0	0	0	0
ONLINE	0	0	0	0	0	0
ONLINE	0	39	0	0	33	32519
ONLINE	0	11	1	0	11	10750
ONLINE	3	11	0	0	12	5154
ONLINE	0	10	9	0	13	7094
ONLINE	0	33	26	0	30	50155
ONLINE	2	37	0	0	37	32419
ONLINE	0	1	0	0	1	4
ONLINE	0	22	19	0	30	24338
ONLINE	2	10	8	0	19	15505
ONLINE	0	1	0	0	1	5



SOLUTION

- ProxySQL "against" ORMMagical
 - "Man-in-the-Middle" Approach
 - Stats on (old/new) Queries
 - The Power of Query Rewrite



SOLUTION

```
mysql> select * from stats_mysql_query_rules
-> where hits > 1000
-> order by hits desc limit 20;
```

rule_id	hits
156	1037641066
321	1020969784
151	832817025
155	828712068
150	828690805
320	594952521
1299	576881271
157	469645197
322	466932589
1154	442881758
1201	429685170
176	405081682
175	402090136
173	382343303
161	361439205
153	278170553
162	248967082
165	171855911
163	170736677
100100	158041540

20 rows in set (0.03 sec)



SOLUTION

```
SELECT COUNT(*) FROM `users` WHERE `users`.`tenant_id` = ?  
AND `users`.`status` = 'active' AND (`users`.`id` IN (  
SELECT `users_applied_tags`.`user_id` FROM  
`users_applied_tags` WHERE `users_applied_tags`.`tag_id` =  
?))
```


digest: 0x285F07C6756FC149

match_pattern: ^SELECT.*tenant_id` = (.*)
AND `users`. *tag_id` = (.*)\$


replace_pattern: SELECT COUNT(*) FROM `users` u
JOIN users_applied_tags uat ON uat.contact_id = u.id
WHERE u.`tenant_id` = \1 AND u.`status` = 'active' AND
uat.`tag_id` = \2

SOLUTION


- ReadySet for Smart Caching
 - Precalculated Results
 - NanoSecond response time
 - Automatic cache invalidation via Change Data Capture
 - Failover mechanism




Smart Caching - ReadySet



Smart Caching - ReadySet



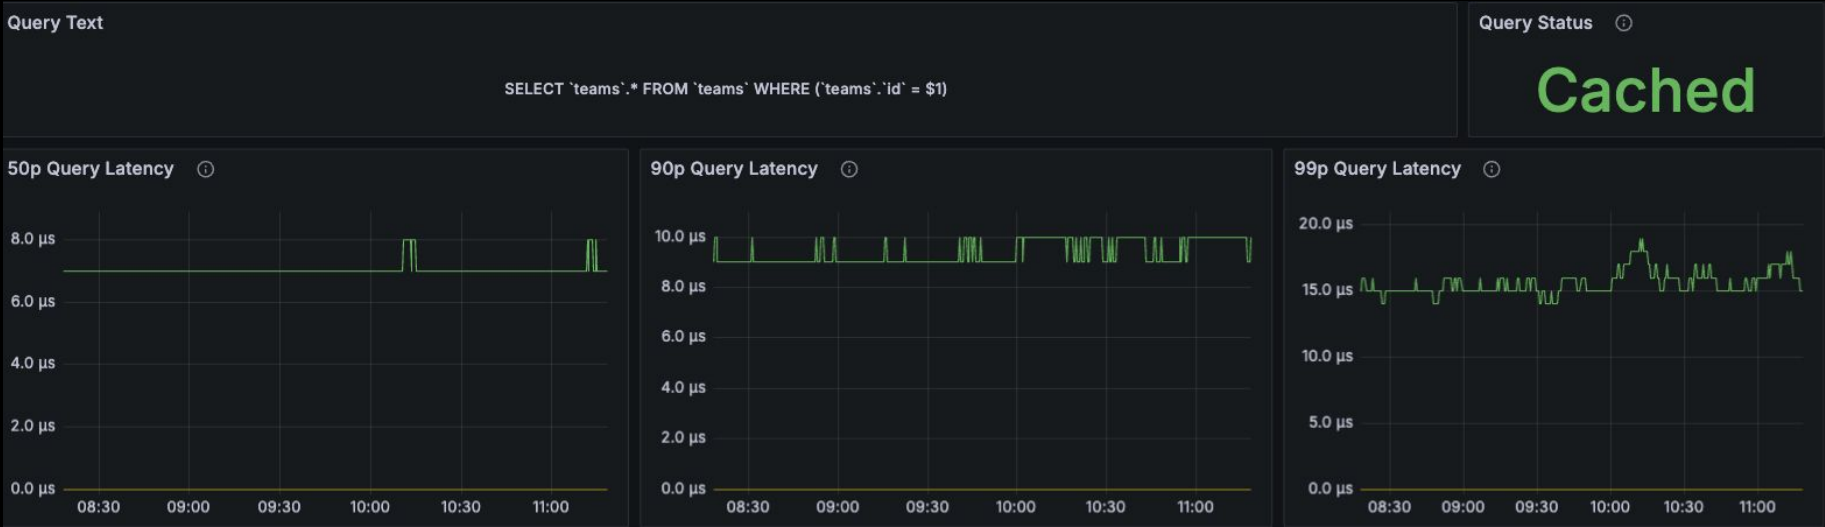
Smart Caching - ReadySet



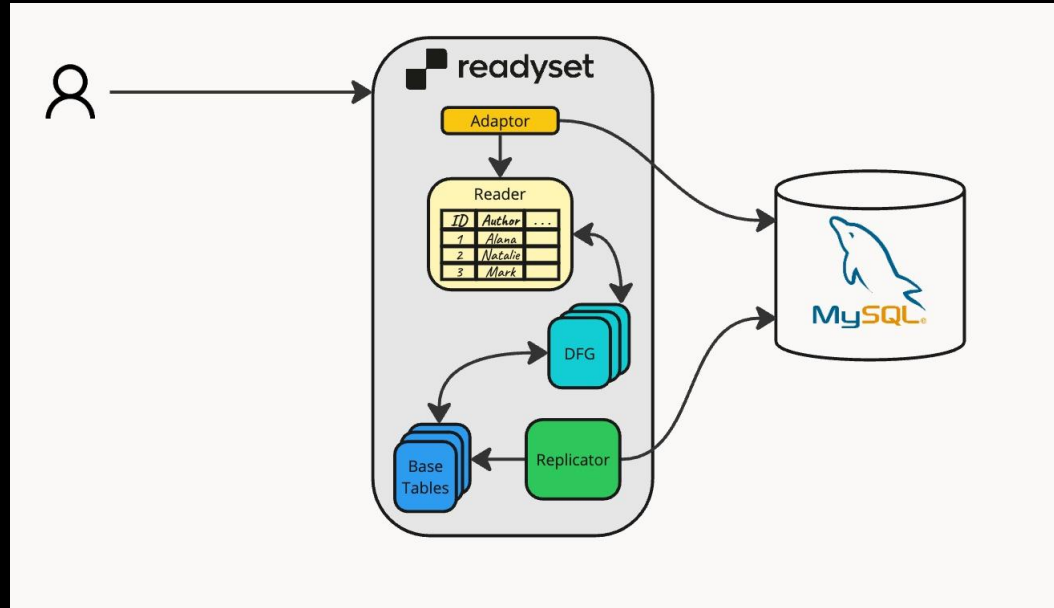
Smart Caching - ReadySet



Smart Caching - ReadySet



Smart Caching - ReadySet



ScaleDB



Questions?



Javier Zon

MySQL Scalability Expert - Taekwondo Instructor

